

Amzur Chatbot Technical Specification

The system uses ChromaDB for vector storage.

Embeddings are generated using OpenAI text-embedding-3-large model.

The LLM model is Google Gemini 2.5 Flash via LiteLLM proxy.

Authentication uses JWT stored in httpOnly cookies.

The backend is built with FastAPI and the frontend uses React 19.

RAG pipeline extracts text from PDFs using PyMuPDF.

Maximum upload size is 100 MB.